

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [3]: df = pd.read_csv(r'C:\Users\sahil\Downloads\WA_Fn-UseC_-Telco-Customer-Churn.csv')
```

```
In [7]: df.head()
```

```
Out[7]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	Mul
--	------------	--------	---------------	---------	------------	--------	--------------	-----

0	7590-VHVEG	Female	0	Yes	No	1	No	
---	------------	--------	---	-----	----	---	----	--

1	5575-GNVDE	Male	0	No	No	34	Yes	
---	------------	------	---	----	----	----	-----	--

2	3668-QPYBK	Male	0	No	No	2	Yes	
---	------------	------	---	----	----	---	-----	--

3	7795-CFOCW	Male	0	No	No	45	No	
---	------------	------	---	----	----	----	----	--

4	9237-HQITU	Female	0	No	No	2	Yes	
---	------------	--------	---	----	----	---	-----	--

5 rows × 21 columns



```
In [9]: df = df.drop("customerID", axis = True)
```

```
In [13]: df.head()
```

```
Out[13]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	In
--	--------	---------------	---------	------------	--------	--------------	---------------	----

0	Female	0	Yes	No	1	No	No phone service	
---	--------	---	-----	----	---	----	------------------	--

1	Male	0	No	No	34	Yes	No	
---	------	---	----	----	----	-----	----	--

2	Male	0	No	No	2	Yes	No	
---	------	---	----	----	---	-----	----	--

3	Male	0	No	No	45	No	No phone service	
---	------	---	----	----	----	----	------------------	--

4	Female	0	No	No	2	Yes	No	
---	--------	---	----	----	---	-----	----	--



```
In [15]: df.describe()
```

Out[15]:

	SeniorCitizen	tenure	MonthlyCharges
<b>count</b>	7043.000000	7043.000000	7043.000000
<b>mean</b>	0.162147	32.371149	64.761692
<b>std</b>	0.368612	24.559481	30.090047
<b>min</b>	0.000000	0.000000	18.250000
<b>25%</b>	0.000000	9.000000	35.500000
<b>50%</b>	0.000000	29.000000	70.350000
<b>75%</b>	0.000000	55.000000	89.850000
<b>max</b>	1.000000	72.000000	118.750000

In [19]: `df.dtypes`

```
Out[19]: gender          object
SeniorCitizen      int64
Partner            object
Dependents         object
tenure             int64
PhoneService       object
MultipleLines      object
InternetService    object
OnlineSecurity     object
OnlineBackup       object
DeviceProtection   object
TechSupport        object
StreamingTV        object
StreamingMovies    object
Contract           object
PaperlessBilling   object
PaymentMethod      object
MonthlyCharges     float64
TotalCharges       object
Churn              object
dtype: object
```

In [29]: `df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')`In [31]: `df.head()`

Out[31]:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	In
--	--------	---------------	---------	------------	--------	--------------	---------------	----

0	Female	0	Yes	No	1	No	No phone service	
1	Male	0	No	No	34	Yes	No	
2	Male	0	No	No	2	Yes	No	
3	Male	0	No	No	45	No	No phone service	
4	Female	0	No	No	2	Yes	No	



In [82]:

```
from sklearn.preprocessing import LabelEncoder
binary_cols = ['gender', 'Partner', 'Dependents', 'PhoneService',
               'PaperlessBilling', 'Churn']
df[binary_cols] = df[binary_cols].replace({'Yes': 1, 'No': 0, 'Male': 1, 'Female': 0})

# One-Hot Encoding for remaining categorical features
df = pd.get_dummies(df, drop_first=True)
```

In [84]:

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
num_cols = ['tenure', 'MonthlyCharges', 'TotalCharges']
df[num_cols] = scaler.fit_transform(df[num_cols])
```

In [106...]

```
X = df.drop('Churn', axis = 1)
y = df['Churn']
```

In [126...]

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.2 , random_s
```

In [128...]

```
df.head()
```

Out[128...]

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	PaperlessBilling
--	--------	---------------	---------	------------	--------	--------------	------------------

0	0	0	1	0	-1.277445	0	
1	1	0	0	0	0.066327	1	
2	1	0	0	0	-1.236724	1	
3	1	0	0	0	0.514251	0	
4	0	0	0	0	-1.236724	1	

5 rows × 31 columns



```
In [130... from sklearn import tree
model = tree.DecisionTreeClassifier()
```

```
In [132... model.fit(X_train,y_train)
```

```
Out[132... ▼ DecisionTreeClassifier ⓘ ?
```

```
DecisionTreeClassifier()
```

```
In [138... model.predict(X_test)
```

```
Out[138... array([1, 1, 0, ..., 0, 0, 1], dtype=int64)
```

```
In [166... model.score(X_test,y_test)
```

```
Out[166... 0.7182398864442867
```

```
In [234... from sklearn.ensemble import RandomForestClassifier
```

```
# Create the model
```

```
model = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
In [236... model.fit(X_train, y_train)
```

```
Out[236... ▼ RandomForestClassifier ⓘ ?
```

```
RandomForestClassifier(random_state=42)
```

```
In [238... model_RandomForestClassifier.predict(X_test)
```

```
Out[238... array([1, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [240... model_RandomForestClassifier.score(X_test,y_test)
```

```
Out[240... 0.7906316536550745
```

```
In [242... pip install xgboost
```

Requirement already satisfied: xgboost in c:\users\sahil\anaconda3\lib\site-packages (3.0.2)Note: you may need to restart the kernel to use updated packages.

Requirement already satisfied: numpy in c:\users\sahil\anaconda3\lib\site-packages (from xgboost) (1.26.4)

Requirement already satisfied: scipy in c:\users\sahil\anaconda3\lib\site-packages (from xgboost) (1.13.1)

```
In [243... from xgboost import XGBClassifier
```

```
xgb_model = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random
```

```
In [244... xgb_model.fit(X_train, y_train)
```

C:\Users\sahil\anaconda3\Lib\site-packages\xgboost\training.py:183: UserWarning: [12:06:23] WARNING: C:\actions-runner\\_work\xgboost\xgboost\src\learner.cc:738: Parameters: { "use\_label\_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
```

Out[244...]

**XGBClassifier**

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric='logloss',
               feature_types=None, feature_weights=None, gamma=None,
               grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
```

In [246...]

```
y_pred = xgb_model.predict(X_test)
y_pred
```

Out[246...]

```
array([1, 0, 0, ..., 0, 0, 1])
```

In [247...]

```
xgb_model.score(X_train, y_train)
```

Out[247...]

```
0.9391196308129216
```

In [232...]

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt


print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Oranges', xticklabels=['No Churn', 'Churn'], yticklabels=['Actual No Churn', 'Actual Churn'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('XGBoost Confusion Matrix')
plt.show()
```

Accuracy: 0.7842441447835344

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.89	0.86	1036
1	0.61	0.50	0.55	373
accuracy			0.78	1409
macro avg	0.72	0.69	0.70	1409
weighted avg	0.77	0.78	0.78	1409



```
In [260...] from sklearn.linear_model import LogisticRegression
```

```
In [262...] model_linear = LogisticRegression()
```

```
In [266...] model_linear = LogisticRegression(max_iter=1000, random_state=42)
```

```
In [274...] from sklearn.impute import SimpleImputer

imputer = SimpleImputer(strategy='mean') # or 'median', 'most_frequent'
X_train = imputer.fit_transform(X_train)
X_test = imputer.transform(X_test)
```

```
In [276...] model_linear.fit(X_train,y_train)
```

```
Out[276...] LogisticRegression
LogisticRegression(max_iter=1000, random_state=42)
```

```
In [280...] model_linear.predict(X_train)
```

```
Out[280...] array([0, 0, 1, ..., 0, 1, 0], dtype=int64)
```

```
In [284...] model_linear.score(X_train,y_train)
```

```
Out[284...] 0.80386936457224
```

```
In [ ]:
```